

Topological Analysis of pLDDT Profiles: Reproducing and Extending Cazals & Sarti (2025)

Algorithms in Structural Bioinformatics — Q4 Project

Giorgos Boulogeorgos

Supervisor: Professor I. Emiris

Co-advisor: P. Rigas

May 2026

Outline

- 1 Motivation
- 2 Mathematical Framework
- 3 Pipeline and Algorithm
- 4 Arity Extension ($Q1 \times Q4$)
- 5 Validation
- 6 Proteome-Scale Results
- 7 Discussion
- 8 Conclusions


What is pLDDT and why does its *structure* matter?

pLDDT (predicted Local Distance Difference Test)

- Per-residue confidence score, $s_i \in [0, 100]$
- Stored in the B-factor column of AlphaFold PDB files
- $s_i \gtrsim 70$: well-folded region
- $s_i \lesssim 50$: intrinsically disordered

Key observation

- Most analyses use pLDDT as a *threshold* (keep/discard)
- **Cazals & Sarti (2025)** treat the full pLDDT *profile along the chain* as a topological object
- This encodes multi-domain and disordered-linker structure



Project goal

Reproduce Q4 pLDDT fragmentation on 23 391 *H. sapiens* fragments; extend with a Q1 arity cross-correlation.

Path graph and scalar filtration


Model: protein chain $\rightarrow$ **path graph** $G_n = (V, E)$

- Vertices $V = \{1, \dots, n\}$: residues
- Edges $E = \{(i, i+1)\}$: peptide bonds
- Scalar function: $u_i = -s_i$ (negative pLDDT)

Filtration: insert vertices in increasing u order (= **decreasing pLDDT** order)

$$\emptyset \subseteq \mathcal{G}_{u_1} \subseteq \dots \subseteq \mathcal{G}_{u_n} = G_n$$

An edge $(i, i+1)$ is **activated** the moment *both* endpoints are inserted.



Numbers = pLDDT. Insertion order = by decreasing pLDDT.

Connected components, Elder Rule, and Persistence Diagram

$N_{cc}(u)$: number of connected components at level u

- New vertex inserted $\Rightarrow N_{cc} + 1$
- Edge joins two components $\Rightarrow N_{cc} - 1$

Elder Rule: when two components merge, the one born at *higher* pLDDT (older) survives. The younger one **dies**, producing a pair:

$$(b_i, d_i) \quad b_i \geq d_i \geq 0$$

Persistence: $p_i = b_i - d_i$

Accretion: $p_i = 0$ (same-level merge)

Toy example: pLDDT = [90, 10, 80, 50, 70]

Event	N_{cc}	PD pair
Insert $i = 1$ (90)	1	—
Insert $i = 3$ (80)	2	—
Insert $i = 5$ (70)	3	—
Insert $i = 4$ (50), edge (4, 5)	3	(70, 50)
edge (3, 4)	2	(80, 50)
Insert $i = 2$ (10), edge (1, 2)	2	(10, 10) [†]
edge (2, 3)	1	(10, 10) [†]

[†]Accretions — zero persistence

Positive persistence pairs: (70, 50) and (80, 50)

Persistence Diagram (PD)

Collection of all $m = n - 1$ pairs $\mathcal{P} = \{(b_i, d_i)\}_{i=1}^{n-1}$

Five topological statistics

Let $m = n-1$, $m' =$ positive-persistence pairs,
 $\mathcal{P}_D =$ set of distinct persistence values.

$$① f_{cp}^+ = \frac{m'}{n-1}$$

Fraction of heterogeneous transitions

$$② \bar{p} = \frac{1}{m} \sum_i (b_i - d_i)$$

Mean persistence lifetime

$$③ H_p = \frac{-\sum_v \mathbb{P}[v] \ln \mathbb{P}[v]}{\ln |\mathcal{P}_D|}$$

Normalised Shannon entropy of persistence


$H_p \in [0, 1]$; high = diverse pLDDT profile

$$④ N_{cc}^{max} = \max_u N_{cc}(u)$$

Peak simultaneous fragmentation


$$⑤ \text{PLM}(t_p): \text{persistent local maxima of } N_{cc}$$

Number of significant domain peaks



$\text{PLM} \geq 3$ at $t_p = 0.025$ flags a protein as a **multi-domain fragmentation candidate**.

End-to-end pipeline overview



6 parallel workers via multiprocessing.Pool — processes 23 391 proteins in ≈ 30 seconds.

Algorithm: standard sublevel-set filtration

Path-graph pLDDT filtration (paper-faithful)

Input: $(s_1, \dots, s_n)$, threshold t_p **Output:** $\mathcal{P}$, N_{cc} -curve, $\text{PLM}(t_p)$

- 1 $\pi \leftarrow \text{STABLEARGSORT}(-s)$ *decreasing pLDDT; ties by index*
- 2 Init UF: $\text{parent}[i] \leftarrow i$, $\text{birth}[i] \leftarrow -\infty$, $\text{inserted}[i] \leftarrow \perp$; $N_{cc} \leftarrow 0$
- 3 **for each** i in π : $\text{birth}[i] \leftarrow s_i$; $\text{inserted}[i] \leftarrow \top$; $N_{cc} \leftarrow N_{cc} + 1$
 - **for** $j \in \{i-1, i+1\}$ with $\text{inserted}[j]$; let $r_i, r_j = \text{FIND}(i, j)$:
 - $r_i = r_j$: append (s_i, s_j) (accretion)
 - **else (Elder Rule)**: $\text{survivor} = \arg \max \text{birth}$; merge loser $\rightarrow$ survivor; $N_{cc} = 1$; append $(\text{birth}[\text{loser}], s_i)$
 - append (s_i, N_{cc}) to Curve
- 4 $\text{PLM}(t_p) \leftarrow \text{GUDHIPERSISTENTMAXIMA}(\text{Curve}, t_\nu = n \cdot t_p)$
- 5 **return** $\mathcal{P}$, Curve, $\text{PLM}(t_p)$

FIND uses path-halving ($O(\alpha(n)) \approx O(1)$). **GUDHI** step: negate N_{cc} , 1D cubical complex, count 0-dim pairs with persistence $\geq t_\nu$.

PLM via GUDHI — second-layer persistence

The N_{cc} curve is itself a 1D scalar function. Its local maxima correspond to **pLDDT domain peaks**.


How we count significant peaks:

- 1 Negate N_{cc} : local maxima become local *minima*
- 2 Feed $-N_{cc}$ to GUDHI's CubicalComplex
- 3 Each 0-dim pair (b_j, d_j) in the resulting PD has persistence = $N_{cc}[\text{peak}] - N_{cc}[\text{valley}]$
- 4 Count pairs with persistence $\geq t_\nu = n \cdot t_p$

Why two layers of persistence?

First layer: filtration of the *protein graph*.

Second layer: Morse-Smale simplification of the N_{cc} *function* — filters out noise peaks from float-precision pLDDT.



Arity: a 3D packing descriptor

Arity of residue k :

of C_α atoms within $10 \text{ \AA}$ of $C_\alpha(k)$ (excluding self)


Arity signature (a_{25} , a_{75}):

25th and 75th percentiles of per-residue arity

- a_{25} low $\Rightarrow$ flexible/disordered regions dominate
- a_{75} high $\Rightarrow$ well-packed globular regions present
- Multi-domain proteins: *both* — compact domains + disordered linkers

Implementation:

- Extract C_α coordinates (gemmi)
- Build `scipy.cKDTree`
- `query_ball_point(r=10)` — $O(n \log n)$
- Subtract 1 for self-count



Validation: toy example and null model

Toy protein ($n = 5$, pLDDT = [10, 90, 80, 90, 10])

Pair	(b_i, d_i)	Type
(80, 80)	accretion	positive
(90, 80)	$p = 10$	
(10, 10)	accretion	
(10, 10)	accretion	

Hand-computed:

$$H_p = 0.8112, \quad f_{cp}^+ = 0.25$$

Code output (4 d.p.):

$$H_p = 0.8113, \quad f_{cp}^+ = 0.25 \checkmark$$

Null random model ($n = 1,000$, pLDDT $\sim \mathcal{U}[0, 100]$)

Cazals & Sarti (2025) Example 1:

- $f_{cp}^+ \approx 1/3 = 0.333$
- $H_p \approx 0.45$
- $N_{cc}^{max} \approx n/4 = 250$

Our result (raw float pLDDT):

- $f_{cp}^+ = 0.329 \checkmark$
- $H_p = 0.438 \checkmark$
- $N_{cc}^{max} = 252 \checkmark$


Note: despite continuous float inputs, only 330 distinct persistence values arise, keeping $H_p \ll 1$.

Validation: three prototype proteins

UniProt	Type	n	H_p	f_{cp}^+	PLM
<i>Paper targets (Fig. 3)</i>					
P15121	Ordered	316	≈ 0.04	—	—
A0A0G2L439	Disord.	449	≈ 0.27	—	—
Q9VQS4	Mixed	781	≈ 0.28	—	—
<i>Our results (raw float pLDDT)</i>					
P15121	Ordered	316	0.371	0.241	1
A0A0G2L439	Disord.	449	0.466	0.339	1
Q9VQS4	Mixed	781	0.424	0.312	2


Key: relative ordering preserved
Ordered < Mixed < Disordered ✓

Why are our H_p values higher?



More distinct pLDDT levels $\rightarrow$ more distinct persistence values $\rightarrow$ broader $\mathbb{P}[v]$ $\rightarrow$ higher H_p .

Proteome-scale Pearson correlation (Figure 7)



Setup: all 23 391 AlphaFold-DB v4 fragments

No length pre-filter (paper-faithful)

Metric	Value
--------	-------

N	23 391
-----	--------

Pearson r	0.864
-------------	--------------

Paper target	≈ 0.97
--------------	----------------

Gap	-0.106
-----	----------

Both statistics elevated (diagonal shift)

→ slope preserved, variance increased

→ r reduced by float-precision effect

Median $f_{cp}^+ = 0.242$, Median $H_p = 0.362$

Fragmentation candidates and t_p sensitivity

Three-way filter (paper's Fig. 8):


$$n \geq 200, \quad H_p \geq 0.25, \quad \text{PLM}(t_p) \geq 3$$

t_p	Candidates	Paper
0.020	822	—
0.025	183	≈ 86
0.030	34	—

Why 183 instead of 86?


Float pLDDT creates many fine-grained N_{cc} oscillations $\rightarrow$ more peaks
clear the threshold $t_v = 0.025 \cdot n$

34 proteins survive all thresholds
= **robust, threshold-independent core**



3D scatter: length $\times H_p \times f_{cp}^+$. Orange = fragmentation candidates.

Q1 × Q4 cross-correlation



Finding: fragmented proteins ($PLM \geq 3$) cluster in a specific arity bin


Property	Value
Centroid	$(a_{25} = 7, a_{75} = 15)$
Enrichment	$5.92\times$
p -value	1.61×10^{-3}
Test	Fisher exact

Biological interpretation

- $a_{25} \approx 7$: disordered linkers with few contacts
- $a_{75} \approx 15$: well-packed domain cores
- Intermediate arity = mix of compact domains + disordered linkers = **multi-domain hallmark**

This finding is *orthogonal* to pLDDT — derived purely from 3D C_α coordinates.

The float-precision gap — unified explanation



Both gaps share a single falsifiable hypothesis: reproducing the analysis on integer-rounded pLDDT would close them — at the cost of departing from the paper's stated methodology.

AlphaFold-DB v6 replication: gap persists

Question: does v6 close the gap with the paper?

Full pipeline re-run on v6 tarball (5.1 GB, May 2026)
23,586 proteins processed, 0 failures

Metric	v4	v6	Δ
N proteins	23,391	23,586	+195
Pearson r	0.864	0.862	-0.002
Median f_{cp}^+	0.242	0.239	-0.003
Median H_p	0.362	0.359	-0.003
Candidates [†]	183	185	+2

[†] $n \geq 200$, $H_p \geq 0.25$, $\text{PLM}(0.025) \geq 3$

Answer: no.

Both v4 and v6 store pLDDT as **float**
→ neither matches the paper's
integer-quantised data


All statistics change by $< 0.5\%$:

- $\Delta r = -0.002$
- $\Delta \text{candidates} = +2$
- +195 new UniProt entries
(curation updates, not
structural changes)

The float-precision gap is an
AlphaFold-DB property,
not a v4 artifact.

Arity enrichment: biological significance

What the enrichment tells us



Two orthogonal signals agree:

- 1 **pLDDT fragmentation (Q4):**
 $PLM \geq 3$ identifies proteins with multiple confident domains separated by linkers
- 2 **C_{α} arity (Q1):**
Moderate a_{25} (linkers) + moderate a_{75} (domains) $\rightarrow$ same proteins

Convergence of two independent structural descriptors provides mutual validation and strengthens confidence in the topological fragmentation signal.

Summary

What we reproduced

- Full path-graph filtration from scratch (Union-Find + Elder Rule)
- All 5 statistics + PLM via GUDHI
- Applied to all 23 391 H. sapiens AlphaFold-DB v4 fragments
- Pearson $r = 0.864$ vs target 0.97
- 183 fragmentation candidates vs target 86
- Relative ordering of prototypes preserved ✓
- Null model targets matched ✓

Novel extension (Q1 $\times$ Q4)

- Proteome-wide arity signatures
- 5.92-fold enrichment of fragmented proteins at ($a_{25} = 7, a_{75} = 15$)
- $p = 1.61 \times 10^{-3}$ (Fisher exact)
- Orthogonal 3D validation of pLDDT fragmentation signal

Key insight

All discrepancies with the paper trace to a single cause: float vs. integer pLDDT in AlphaFold-DB v4 vs. the paper's data.

Code: Python 3.11 — NumPy · SciPy · gemmi · GUDHI · PyArrow · streaming tar pipeline (30 s for 23 391 proteins)

Thank you

Questions?

Paper: Cazals & Sarti (2025), bioRxiv 2024.11.16.623929

Data: AlphaFold-DB v4 — *H. sapiens* proteome

